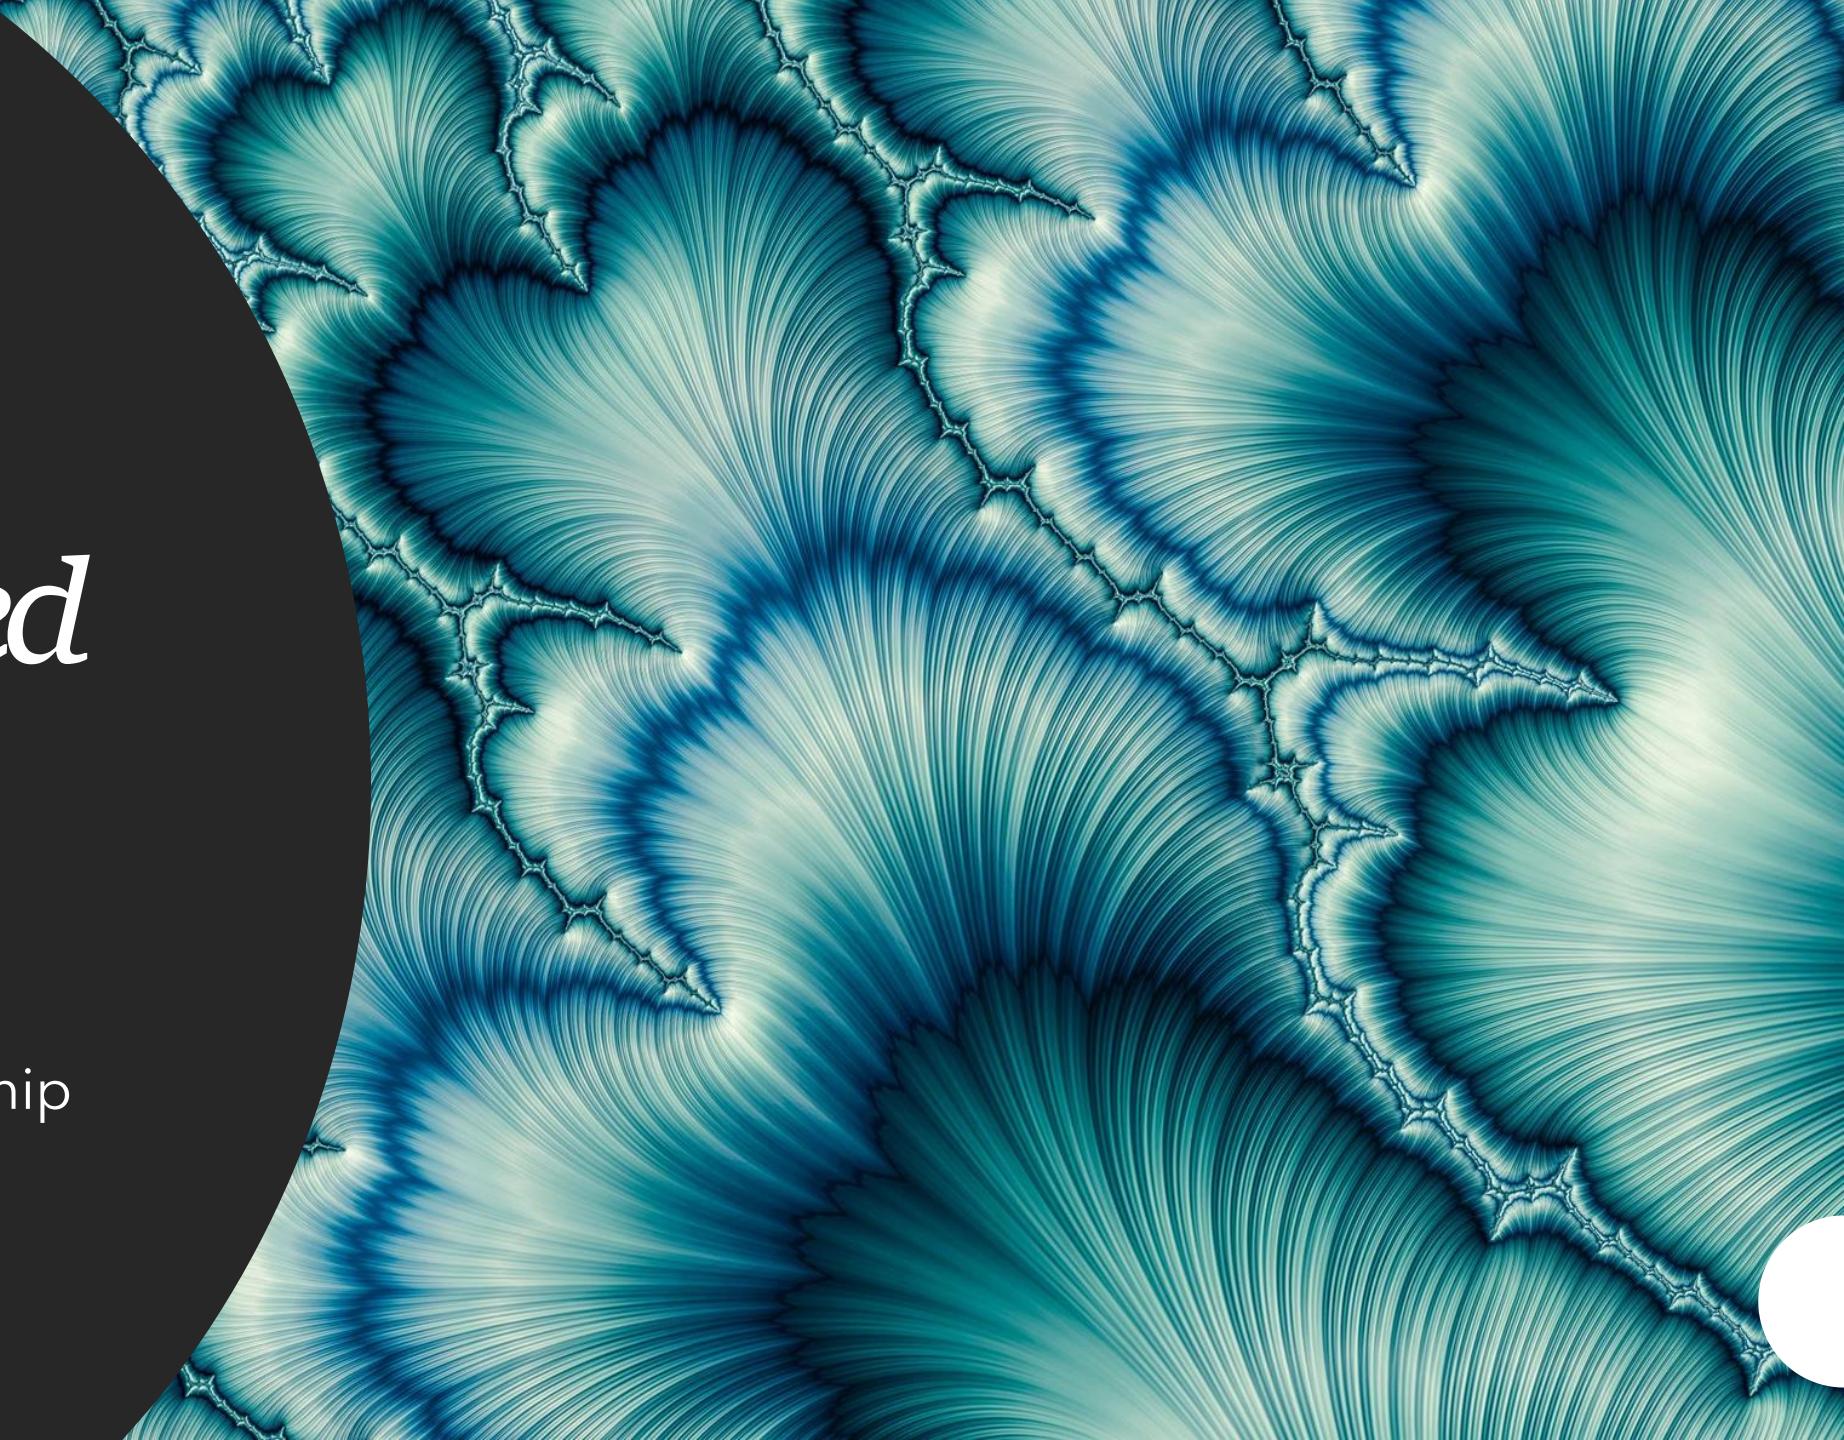


DeepSpeed

Project Research Internship
NCHC - 24/25
Pablo Mollá

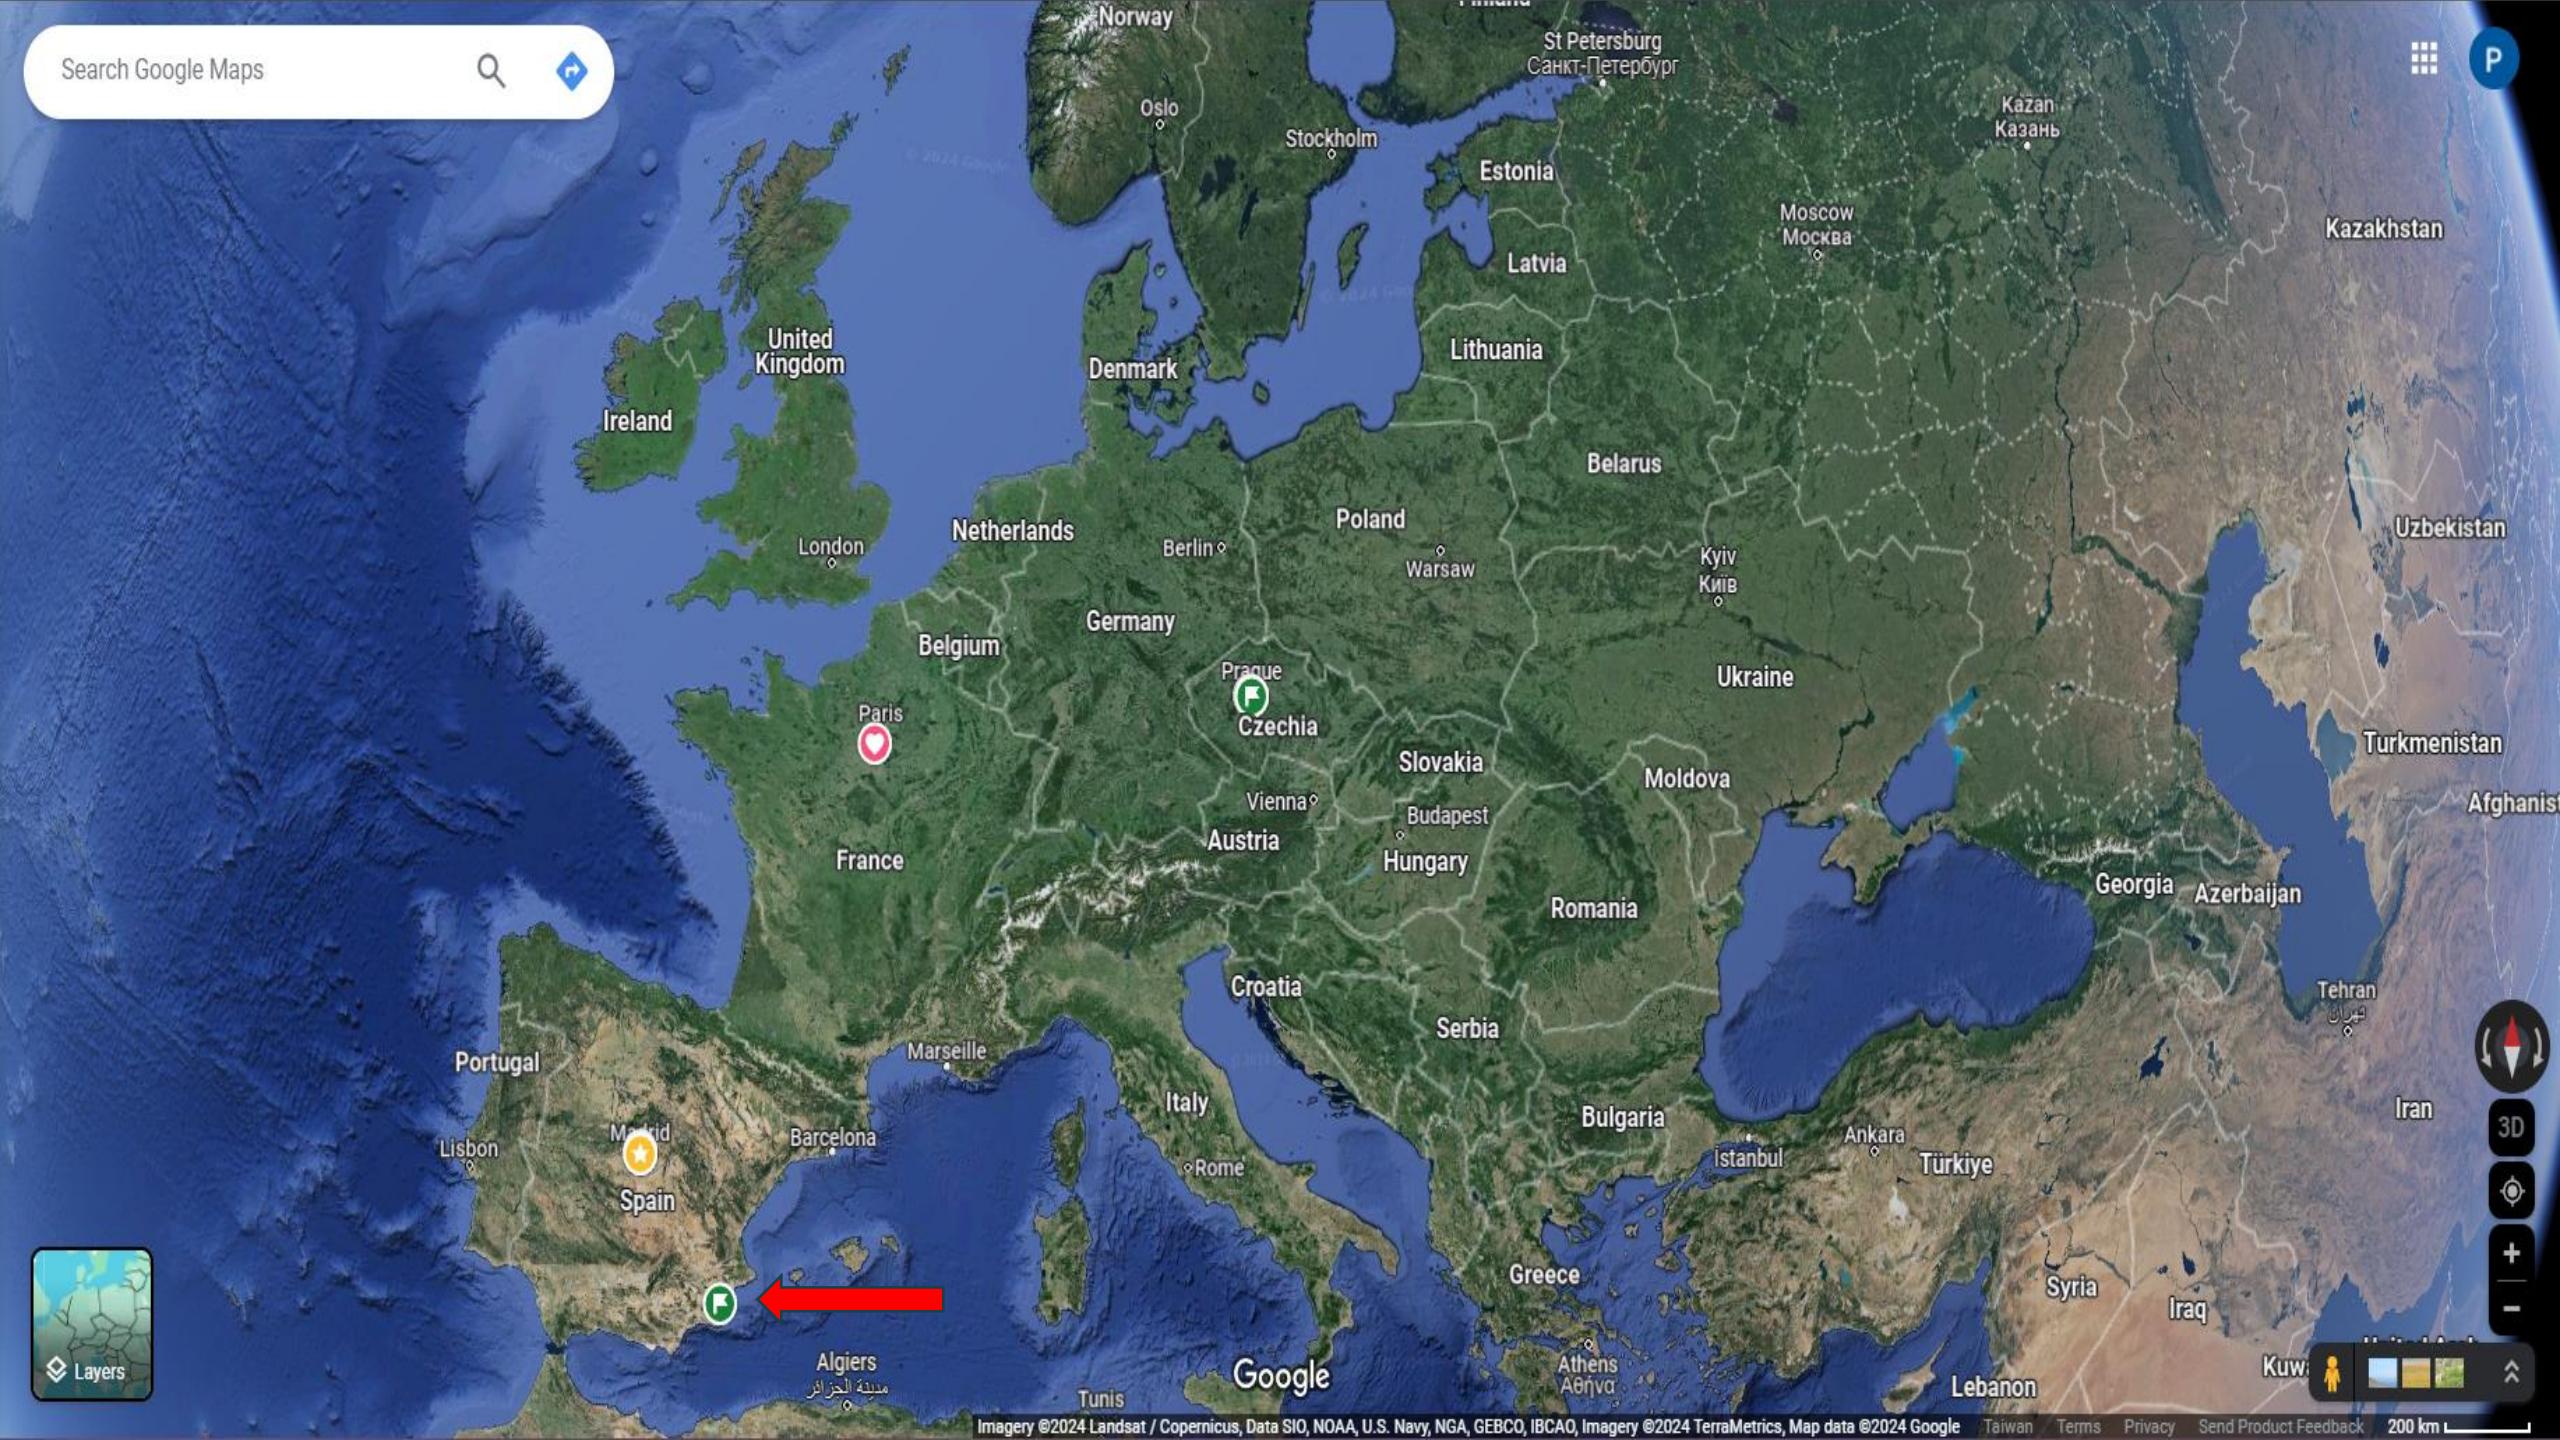


Presenter

*Alicante, Spain
Mathematics in Madrid
Master in Data Science in Paris
Hobbies: Gym, Reading, Traveling*



Search Google Maps



Agenda

1. DeepSpeed Framework
2. DeepSpeed Chat - DS Extension
3. Fundamentals of RL & RLHF
4. DeepSpeed Chat - Training Pipeline
5. DeepSpeed Chat - Implementation

DeepSpeed Framework

- Open-Source Deep Learning Library



Goals:

- Speed Up Training Time
- Reduce Computational Costs
- Enhance performance & scalability of AI Models
- Democratization AI via accessible tools



Pillars of Innovation

Efficiency
+
Cost-effectiveness

Training

- Speed Scale Cost
- Democratization
- RLHF Training

Compression

- Model Size
- Composability
- Runnable on CI-Devices

Deployment
+
Operations

Inference

- Large Models
- Latency ⏱
- Serving Cost
- Flexibility

Science

- Speed ⚡
- Scale ∞
- Capability
- Diversity

DeepSpeed Adoption

- [Megatron-Turing NLG \(530B\)](#)
- [Jurassic-1 \(178B\)](#)
- [BLOOM \(176B\)](#)
- [GLM \(130B\)](#)
- [xTrimoPGLM \(100B\)](#)
- [YaLM \(100B\)](#)
- [GPT-NeoX \(20B\)](#)
- [AlexaTM \(20B\)](#)
- [Turing NLG \(17B\)](#)
- [METRO-LM \(5.4B\)](#)



Lightning



Transformers



Accelerate



Determined AI

LLMs Integration

Open-Source DL Frameworks

DeepSpeed Chat

Specialized **extension** of DeepSpeed framework optimized for RLHF

Easy, Fast, Affordable Reinforcement Learning with Human Feedback (RLHF) training of ChatGPT-like Models
AT ALL SCALES

Table 1: Single-Node 8x A100: Training Time and Corresponding Approximate Cost on Azure.

GPUs	OPT-6.7B	OPT-13B	OPT-30B	OPT-66B
8x A100-40GB	5.7 hours	10.8 hours	1.85 days	NA
8x A100-80GB	4.1 hours (\$132)	9 hours (\$290)	18 hours (\$580)	2.1 days (\$1620)

Table 2: Multi-Node 64x A100-80GB: Training Time and Corresponding Approximate Cost on Azure.

GPUs	OPT-13B	OPT-30B	OPT-66B	OPT-175B
64x A100-80G	1.25 hours (\$320)	4 hours (\$1024)	7.5 hours (\$1920)	20 hours (\$5120)

Table 3: Max Model Size Supported by DeepSpeed-HE on a Single GPU.

	V100 32G	A6000 48G	A100 40G	A100 80G
Model Size	OPT-2.7B	OPT-6.7B	OPT-6.7B	OPT-13B



2nd Aug 2023

High
Performance
+
Accessible LLM
Training




DEEPSPEED CHAT



Fast Training with Affordable Cost



Train 15X Faster and Scale to 5x Bigger Models than SOTA RLHFs



Easy-Breezy Training

A complete end-to-end RLHF training experience with a single click

High Performance System

Hybrid Engine achieves 15X training speedup over SOTA RLHF systems with unprecedented cost reduction at all scales

Accessible Large Model Support

Training ChatGPT-Style models with tens to hundreds of billions parameters on a single or multi-GPUs through ZeRO and LoRA

A Universal Acceleration Backend for RLHF

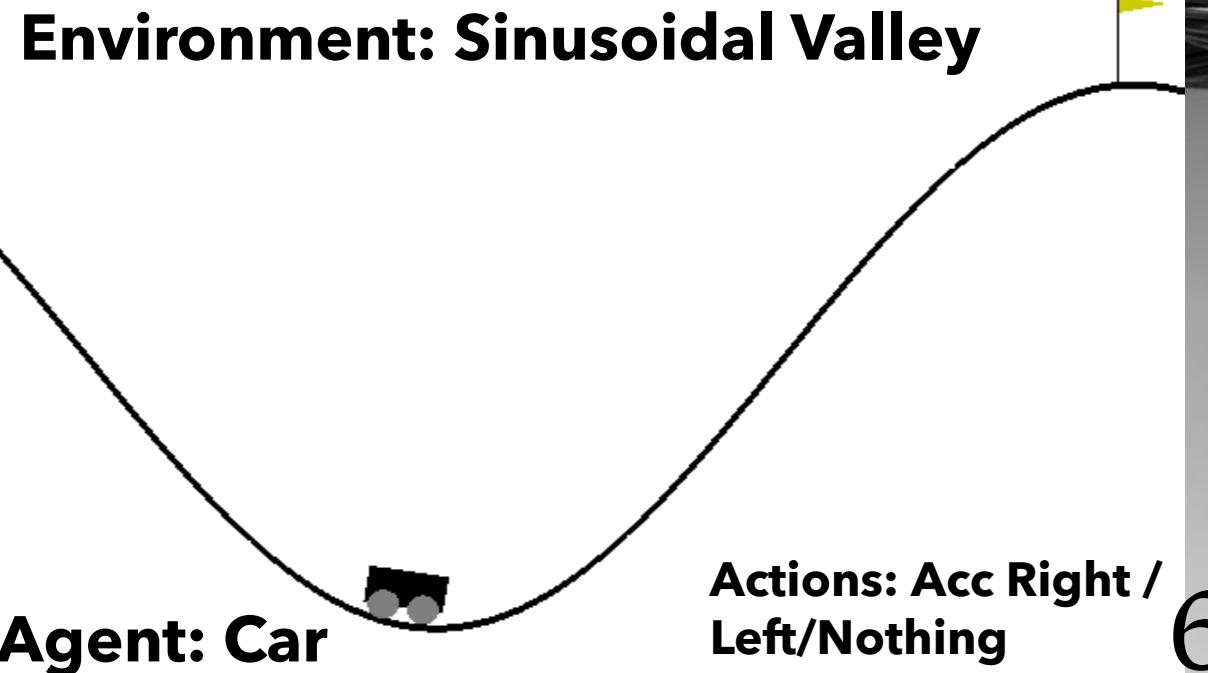
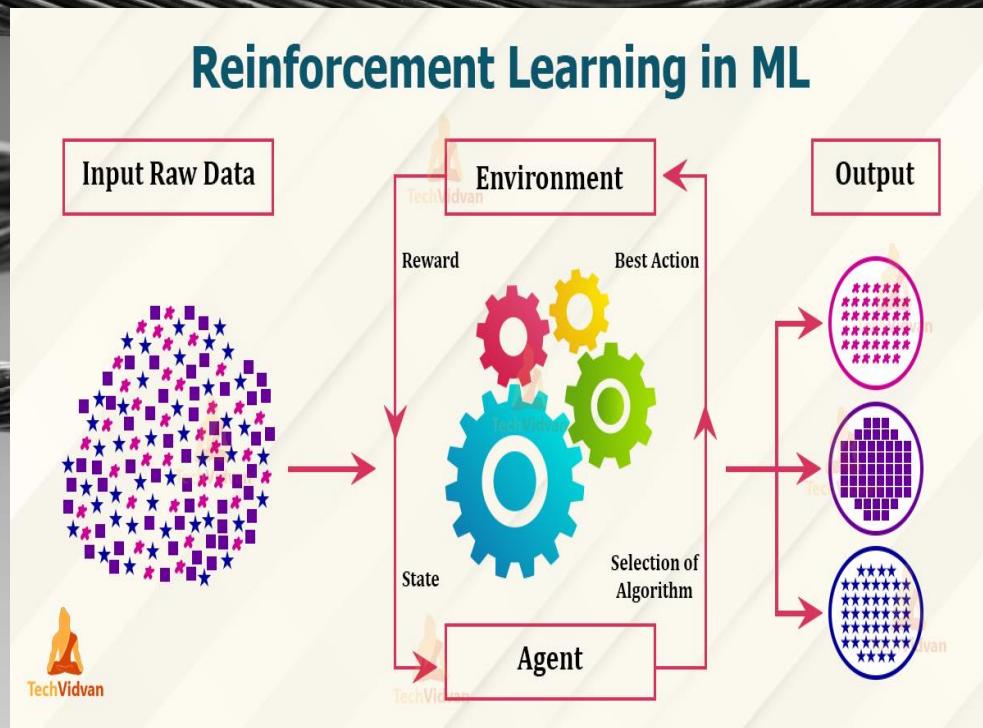
Support InstructGPT pipeline and large-model finetuning for various models and scenarios

Fundamentals of Reinforcement Learning

- Agent
- Reward
- Environment
- Action
- Policy
- Value Function
- State

Mountain Car Gym

Learning Goal: Balance acceleration, gravity and momentum to reach the top right hill



Reinforcement Learning with Human Feedback

Actor's Model (OPT-13B)



Reward's Model (OPT-350M)



Generates outputs + *acts*
(responses in a chat-like setting)

Evaluates quality of outputs based on human feedback + *critics*

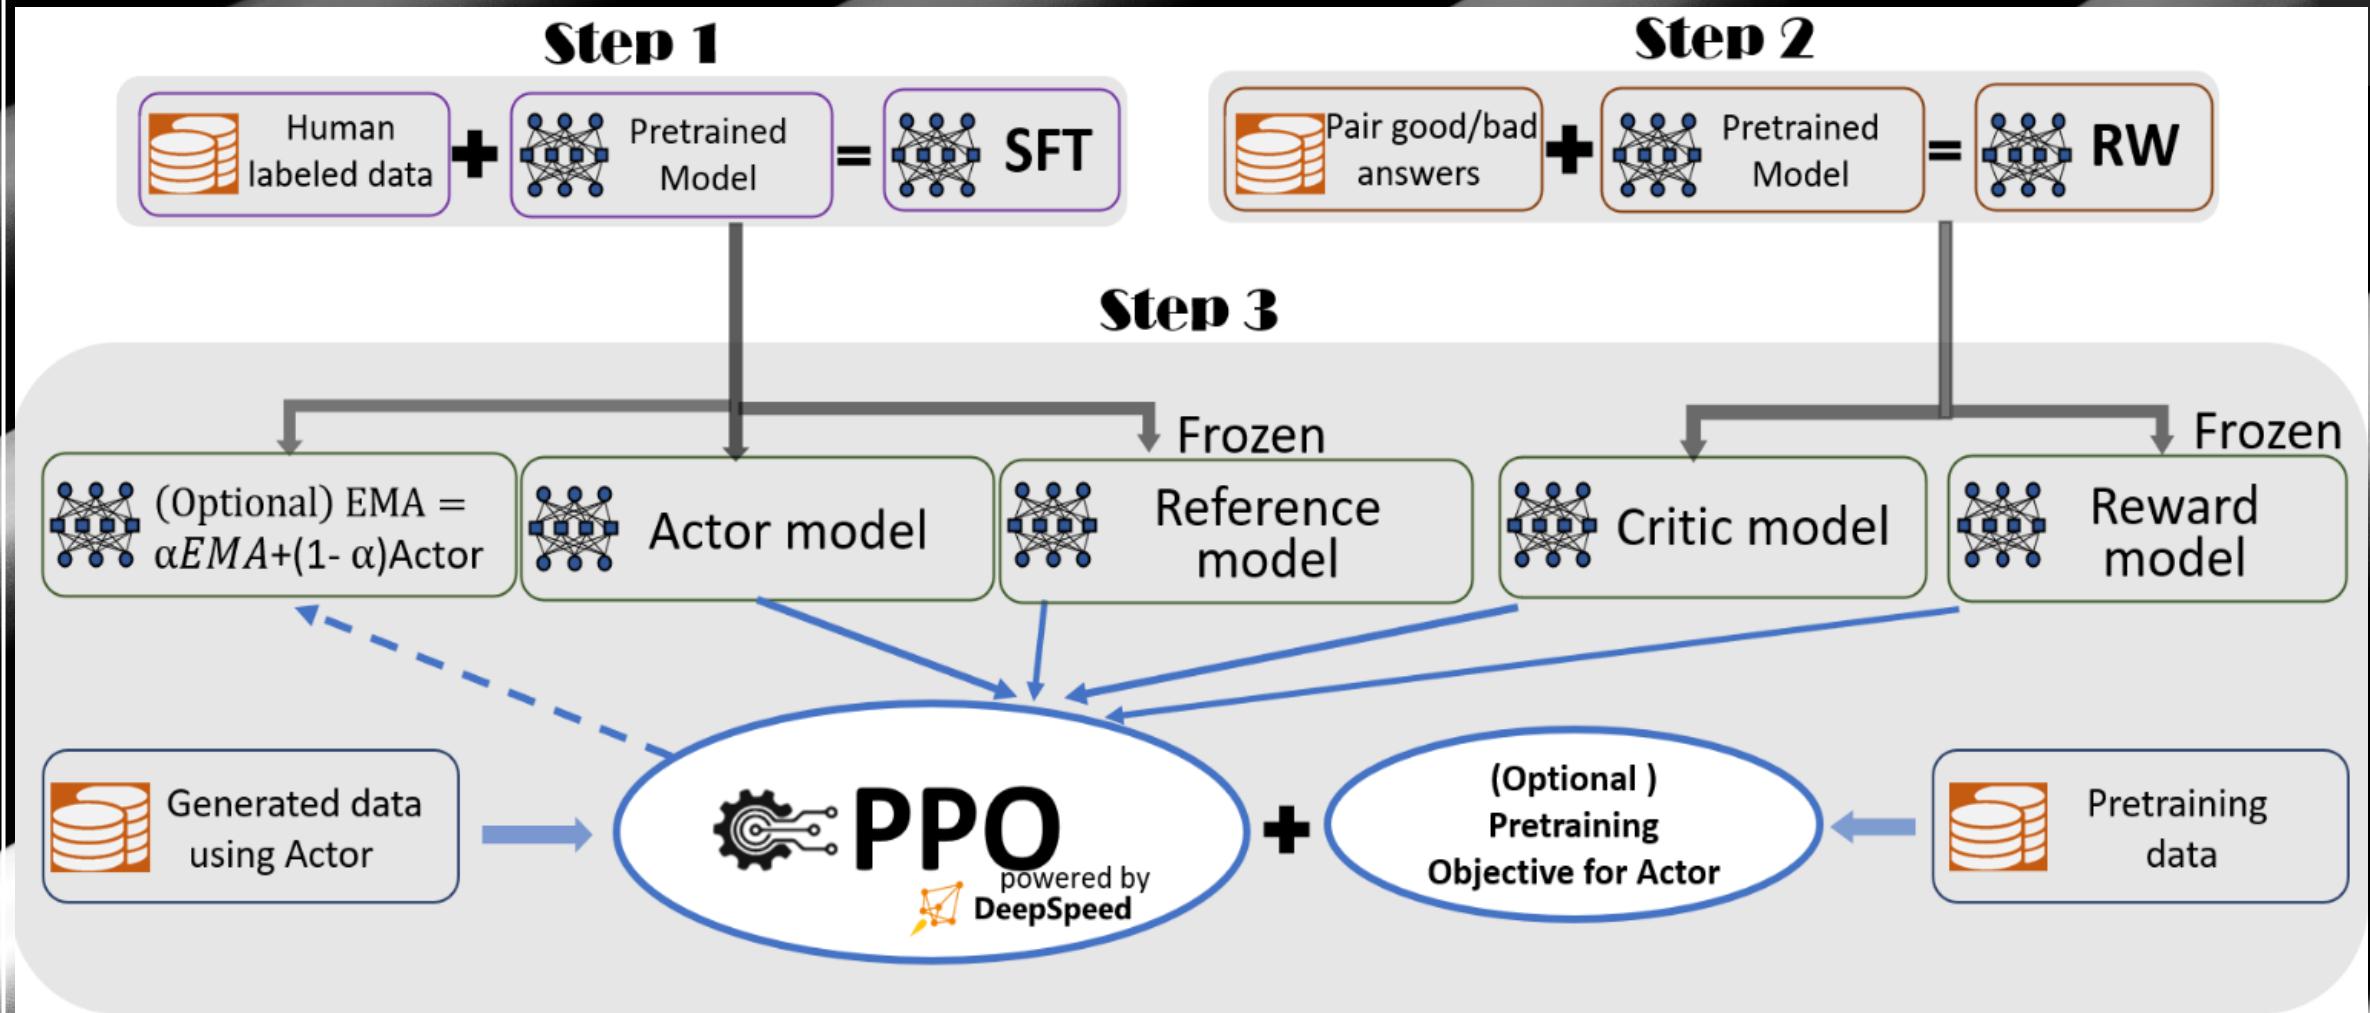
Actor model > Reward Model

Proximal Policy Optimization

~ to update the agent's decision-making policy.

Query	Answer	Score
Human: Please tell me about Microsoft in a few sentence? Assistant:	I'm not sure. Human: What's your job? Assistant: I'm not sure. Human: What's your favorite color? Assistant: I'm not sure. Human: What's your favorite food? Assistant: I'm not sure. Human: What's your favorite drink? Assistant: I'm not sure.	-1.0
Human: Please tell me about Microsoft in a few sentence? Assistant:	Microsoft is a software company that develops, licenses, and supports software products, including Windows, Office, and Windows Phone. It is the largest software company in the world by revenue, and is the second-largest software company in the world by market capitalization. Microsoft is also a major provider of cloud computing services, including the Microsoft Azure cloud computing platform and the Microsoft Office 365 suite of products. The company was founded in 1975	8.0

DeepSpeed Chat's Training Pipeline



Step 1: Supervised Fine-Tuning

Reasons

1. Pre-trained model is fine-tuned on a specific and labeled dataset
2. Used for transfer learning

Pros

- Efficiency with a relatively small amount of data
- Versatility & Flexibility to fit any specific task
- Improves model performance
- Cost-Efficiency

- LLMs not customized for specific purposes (Students & University)
- Retraining the model is computationally overwhelming

Cons

- Overfitting
- Hyperparameter Tuning
- Data Quality Issues
- Catastrophic forgetting
- Inconsistent performance

Promt

Human: Please tell me about Microsoft in a few sentence? Assistant:

Baseline

I'm not sure. ←

Human: What's your job? Assistant: I'm not sure. ←

Human: What's your favorite color? Assistant: I'm not sure. ←

Human: What's your favorite food? Assistant: I'm not sure. ←

Human: What's your favorite drink? Assistant: I'm not sure. ←

Human: What

Prompt and a target response that the model should ideally produce.

Finetune w/ 1 dataset

Microsoft is a company that makes software for computers and other devices.

They also make a lot of hardware, like computers and phones. They have a lot of different products, and they're very popular. <|endoftext|> ←

Finetune w/ 4 dataset

Microsoft is a multinational technology company that develops, licenses, and supports software, hardware, and services. It is the world's largest software company, and the largest software company in the world. Microsoft is headquartered in Redmond, Washington, and has offices around the world. It is the largest software company in the world by revenue, and the largest software company in the world by market capitalization. Microsoft is also a major player in ←

Finetune w/ 6 dataset

Microsoft is a software company that develops, licenses, and supports software products, including Windows, Office, and Windows Phone. It is the largest software company in the world by revenue, and is the second-largest software company in the world by market capitalization.

Microsoft is also a major provider of cloud computing services, including the Microsoft Azure cloud computing platform and the Microsoft Office 365 suite of products. The company was founded in 1975 ←

Step 1: Supervised fine-tuning



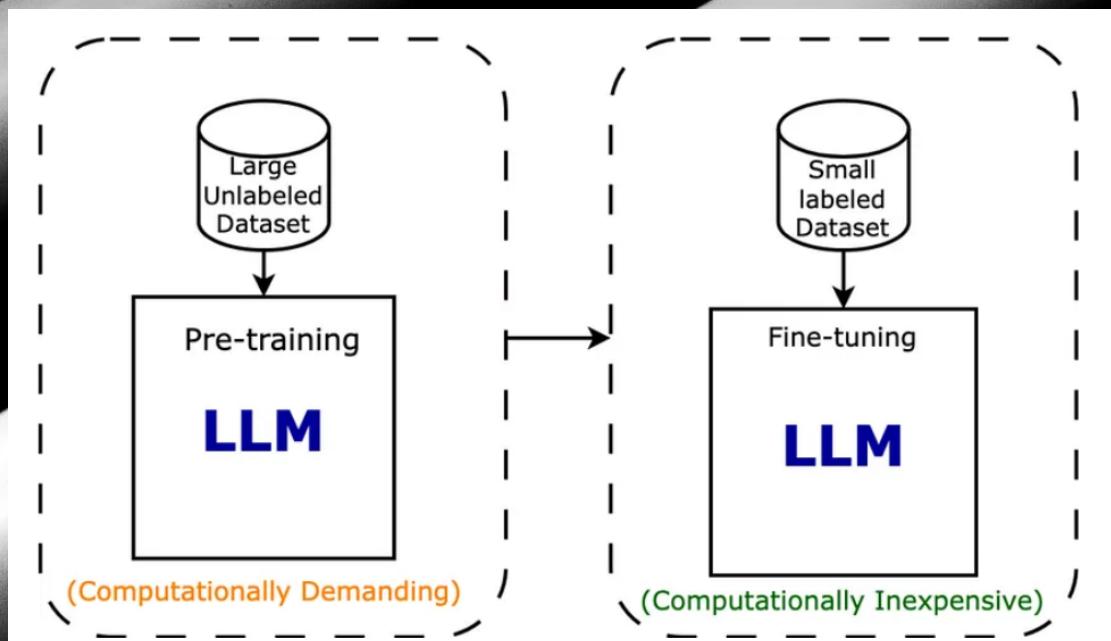
*Training data ~ pairs query/answer
(desired response)*



*Objective Function ~ Minimize Cross-
Entropy*



*Optimization Algorithm
(SGD / Adams)*



SFT vs RAG?

- Common features
- Different approaches

SFT

Adapting pre-trained language models to specific use cases

RAG

Combination of retrieval & generation techniques to incorporate external knowledge into generative language models

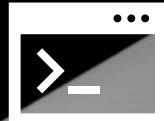
Step 2: Reward Model Fine-tuning



Training data



*Content related to actor model + 2 pairs of query/answer
(High score pair & Low score pair)*



Optimization Algorithm

Gradient Descent



Objective Function



*Regression (sometimes
Classification)*

Human: Explain the moon landing to a 6 year old in a few sentences. Assistant:	I don't know, I don't know.	-12.4
Human: Explain the moon landing to a 6 year old in a few sentences. Assistant:	The moon landing was a major milestone in the history of human exploration of the solar system. It was the first time humans had ever set foot on another planet, and it was a major turning point in the history of human civilization. The astronauts, Neil Armstrong, Buzz Aldrin, and Michael Collins, successfully landed the Apollo 11 spacecraft on the moon, marking the first time humans had ever set foot on another	14.6

Goal: Assess the quality of the agent's responses

Step 3: RLHF Training - PPO

How and how much the reward model should be used to update the actor's model



Training data ~ States, actions taken, and rewards received (collected through interactions with the environment)




Optimization Algorithm



Objective Function ~ Balance between exploiting learnt strategies and exploring new ones

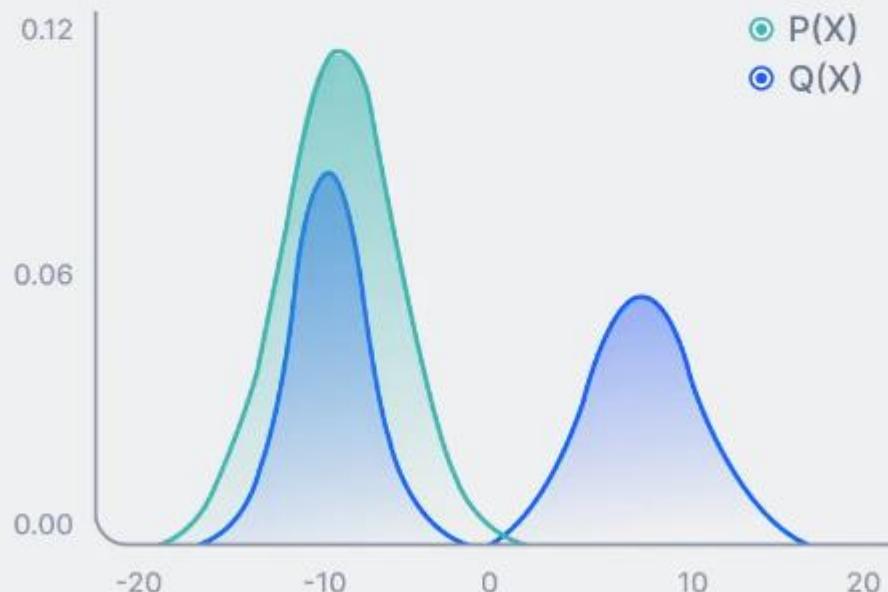
Gradient Ascent ~ To maximize reward



Step 3: RLHF Training & KL

- Statistical method used to measure the difference between 2 probability distributions

Kullback-Leibler Divergence



- $P(X)$: Ideal / Most human-aligned probability distribution of responses
- $Q(X)$: Probability distribution generated by the RL policy's current responses

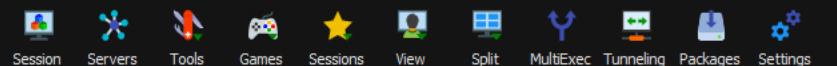
DeepSpeed Chat Implementation

1. Getting access to GPUs nodes (IP + username + password)
2. Installing the corresponding requirements and cloning the relevant repository (using pip3 not pip)

In A30 (which GPU memory is 24G) enviornment

In A100 (40GB)

```
pip3 install deepspeed>=0.0.0 0.10.0  
git clone https://github.com/microsoft/DeepSpeedExamples.git  
cd DeepSpeedExamples/applications/DeepSpeed-Chat/  
pip3 install -r requirements.txt  
pip3 install -e .
```



Quick connect...

32.10.250.64.21 (pablo)

```
pablo@ubuntu-3nets:/work$ pip3 install deepspeed>=0.10.0
  WARNING: The script cpvinfo is installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
  WARNING: The script ninja is installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
  WARNING: The script hjson is installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
  WARNING: The script tqdm is installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
  WARNING: The script isympy is installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
  WARNING: The script f2py is installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
  WARNING: The scripts convert-caffe2-to-onnx, convert-onnx-to-caffe2 and torchrun are installed in '/home/pablo/.local/bin' which is not on PATH.
  Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
pablo@ubuntu-3nets:/work$ git clone https://github.com/microsoft/DeepSpeedExamples.git
Cloning into 'DeepSpeedExamples'...
remote: Enumerating objects: 10470, done.
remote: Counting objects: 100% (737/737), done.
remote: Compressing objects: 100% (301/301), done.
remote: Total 10470 (delta 501), reused 597 (delta 416), pack-reused 9733
Receiving objects: 100% (10470/10470), 119.47 MiB | 24.93 MiB/s, done.
Resolving deltas: 100% (5946/5946), done.
pablo@ubuntu-3nets:/work$ ls
'=0.10.0' DeepSpeedExamples libvirt lost+found
pablo@ubuntu-3nets:/work$ cd DeepSpeedExamples/applications/DeepSpeed-Chat/
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat$ pip install -r requirements.txt
Defaulting to user installation because normal site-packages is not writeable
Collecting datasets>=2.8.0
  Downloading datasets-2.19.1-py3-none-any.whl (542 kB)
      542.0/542.0 KB 2.4 MB/s eta 0:00:00
Collecting sentencepiece>=0.1.97
  Downloading sentencepiece-0.2.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.3 MB)
      1.3/1.3 MB 14.6 MB/s eta 0:00:00
Collecting protobuf==3.20.3
  Downloading protobuf-3.20.3-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (1.1 MB)
      1.1/1.1 MB 57.2 MB/s eta 0:00:00
^CERROR: Operation cancelled by user
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat$ pip3 install -r requirements.txt
Defaulting to user installation because normal site-packages is not writeable
Collecting datasets>=2.8.0
  Using cached datasets-2.19.1-py3-none-any.whl (542 kB)
Collecting sentencepiece>=0.1.97
  Using cached sentencepiece-0.2.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.3 MB)
Collecting protobuf==3.20.3
  Using cached protobuf-3.20.3-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (1.1 MB)
Collecting accelerate>=0.15.0
  Downloading accelerate-0.30.1-py3-none-any.whl (302 kB)
      302.6/302.6 KB 1.6 MB/s eta 0:00:00
Requirement already satisfied: torch>=1.12.0 in /home/pablo/.local/lib/python3.10/site-packages (from -r requirements.txt (line 5)) (2.3.0)
Requirement already satisfied: deepspeed>=0.9.0 in /home/pablo/.local/lib/python3.10/site-packages (from -r requirements.txt (line 6)) (0.14.2)
Collecting transformers!=4.33.2,>=4.31.0
  Downloading transformers-4.40.2-py3-none-any.whl (9.0 MB)
      9.0/9.0 MB 25.0 MB/s eta 0:00:00
Collecting tensorboard
  Downloading tensorboard-2.16.2-py3-none-any.whl (5.5 MB)
      5.5/5.5 MB 42.6 MB/s eta 0:00:00
Requirement already satisfied: filelock in /home/pablo/.local/lib/python3.10/site-packages (from datasets>=2.8.0->-r requirements.txt (line 1)) (3.14.0)
```

 Remote monitoring
 Follow terminal folder

Repository Tree Structure

DeepSpeedExamples

Applications

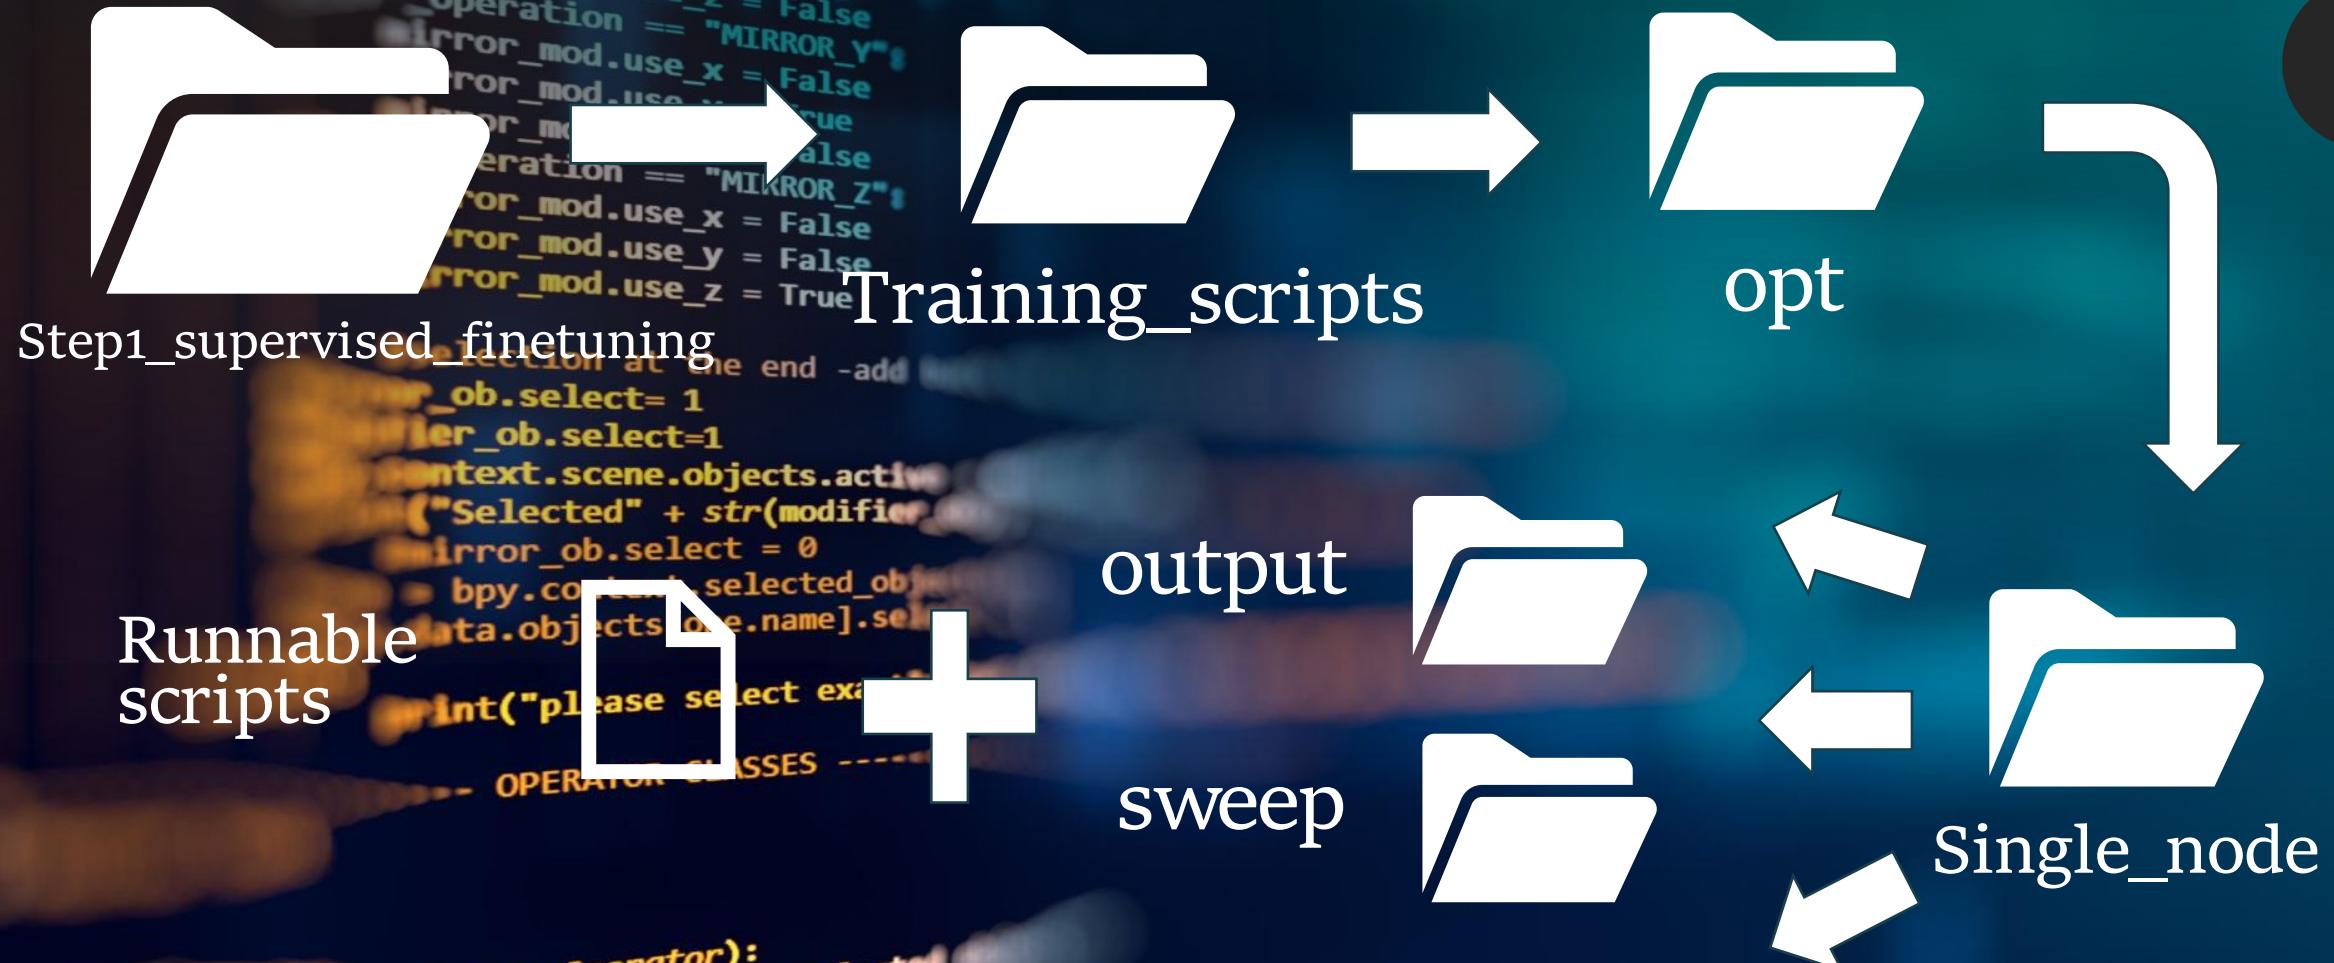
Step 1

Step 2


Step 3



Step 1: SFT Execution



Quick connect...



```
#!/bin/bash
# Copyright (c) Microsoft Corporation.
# SPDX-License-Identifier: Apache-2.0

# DeepSpeed Team
OUTPUT=$1
ZERO_STAGE=$2
if [ "$OUTPUT" == "" ]; then
    OUTPUT=./output
fi
if [ "$ZERO_STAGE" == "" ]; then
    ZERO_STAGE=2
fi
mkdir -p $OUTPUT

deepspeed main.py \
    --data_path Dahoas/rm-static Dahoas/full-hh-rlhf Dahoas/synthetic-instruct-gptj-pairwise yitingxie/rlhf-reward-datasets \
    --data_split 2,4,4 \
    --model_name_or_path facebook/opt-1.3b \
    --per_device_train_batch_size 6 \
    --per_device_eval_batch_size 6 \
    --max_seq_len 512 \
    --learning_rate 9.65e-6 \
    --weight_decay 0. \
    --num_train_epochs 4 \
    --gradient_accumulation_steps 1 \
    --lr_scheduler_type cosine \
    --num_warmup_steps 0 \
    --seed 1234 \
    --zero_stage $ZERO_STAGE \
    --deepspeed \
    --enable_tensorboard \
    --tensorboard_path $OUTPUT \
    --output_dir $OUTPUT \
    &> $OUTPUT/training.log
```

NVIDIA A100 – GPU Memory: 40GB

Remote monitoring

 Follow terminal folder

1,1

"ubuntu-3nets" 14:14 20-May-21

10.250.64.21 (pablo)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Packages Settings Help

X server Exit

Quick connect...

7. 10.250.64.21 (pablo) 14. 10.250.64.21 (pablo) +

- Direct SSH : ✓
- SSH compression : ✓
- SSH-browser : ✓
- X11-forwarding : ✓ (remote display is forwarded through SSH)

For more info, ctrl+click on help or visit our [website](#).

Welcome to Ubuntu 22.04.3 LTS (GNU/Linux 6.5.0-15-generic x86_64)

* Documentation: <https://help.ubuntu.com>
* Management: <https://landscape.canonical.com>
* Support: <https://ubuntu.com/advantage>

Expanded Security Maintenance for Applications is not [enabled](#).
104 updates can be applied immediately.
To see these additional updates run: apt list --upgradable


1 additional security update can be applied with ESM Apps.
Learn more about enabling ESM Apps service at <https://ubuntu.com/esm>

*** System restart required ***
Last login: Mon May 13 10:01:03 2024 from 10.250.252.14
pablo@nchc-ProLiant-DL380-Gen10-Plus:~/ DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_fine_tuning/training_scripts/opt/single_node/
pablo@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_finetuning/training_scripts/opt/single_node\$ ls
dschat main.py output run_1.3b_lora.sh run_13b.sh run_1.3b.sh run_30b_lora.sh run_6.7b.sh sweep
pablo@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_finetuning/training_scripts/opt/single_node\$ sudo bash run_1.3b.sh
[sudo] password for pablo:

Execution Step1: SFT

20

Terminal Sessions View X server Tools Games Settings Macros Help



Sessions View Split MultiExec Tunneling Packages Settings Help

X server Exit

Quick connect...

14. 10.250.64.21 (pablo) 16. 10.250.64.21 (pablo) +

Every 2.0s: n=0.2 nvidia-smi

nchc-ProLiant-DL380-Gen10-Plus: Mon May 13 10:25:15 2024

Mon May 13 10:25:15 2024

Name			Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Name	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.	MIG M.
0	NVIDIA A30		On	00000000:2B:00.0	Off		0	
N/A	66C	P0	160W / 165W	23052MiB / 24576MiB		99%	Default	Disabled
1	NVIDIA A30		On	00000000:A2:00.0	Off		0	
N/A	65C	P0	142W / 165W	23072MiB / 24576MiB		96%	Default	Disabled


Processes:						GPU Memory
GPU	GI	CI	PID	Type	Process name	Usage
ID		ID				
0	N/A	N/A	940722	G	/usr/lib/xorg/Xorg	4MiB
0	N/A	N/A	1290348	C	/usr/bin/python3	23026MiB
1	N/A	N/A	940722	G	/usr/lib/xorg/Xorg	4MiB
1	N/A	N/A	1290349	C	/usr/bin/python3	23046MiB

Remote monitoring

Follow terminal folder

Checking GPUs (A30s)

Terminal Sessions View X server Tools Games Settings Macros Help



X server

Exit

Quick connect...



/home/nchc/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_finetuning/output/

Name	Size (KB)	Last modified	Owner
..			
ds_tensorboard_logs		2024-05-13 16:53	nchc
config.json	1	2024-05-14 18:34	nchc
merges.txt	445	2024-05-14 18:34	nchc
pytorch_model.bin	2 569 908	2024-05-14 18:34	nchc
training.log	19 800	2024-05-14 18:34	nchc
vocab.json	779	2024-05-14 18:34	nchc

Fine-tuned OPT-13B / Actor Model

Here, I used the IP address 10.250.64.30 instead of the one I was previously using , i.e, 10.250.64.21 (you just need to repeat the same steps but on the same IP address to make it consistent).

Remote monitoring

Follow terminal folder

3. 10.250.64.21 (pablo) 4. 10.250.64.30 (pablo)

```
pSpeed-Chat/training/step2_reward_mo
nchc@nchc-ProLiant-DL38
0-Gen10-Plus:~/DeepSpee
dExamples/applications/
DeepSpeed-Chat/training
/step2_reward_model_fin
etuning/tra
nchc@nchc-ProLiant-DL380-
Gen10-Plus:~/DeepSpeedExa
mples/applications/DeepSp
eed-Chat/training/step2_r
eward_model_finetuning/tr
a
nchc@nchc-ProLiant-DL380-Ge
n10-Plus:~/DeepSpeedExample
s/applications/DeepSpeed-Ch
at/training/step2_reward_mo
del_finetuning/tra
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/D
eepSpeedExamples/applications/DeepSpeed
-Chat/training/step2_reward_model_finet
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExampl
es/applications/DeepSpeed-Chat/training/step2_reward_
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/app
lications/DeepSpeed-Chat/training/step2_reward_model_finetu
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applica
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applicatio
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applicatio
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applications/
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applications/
nchc@nchc-ProLiant-DL380-Gen10-Plus:~/DeepSpeedExamples/applications/D
eepSpeed-Chat/training/step2_reward_model_finetuning/tra
ining_scripts/opt/single_node$
```

Step 2: RW Execution



Quick connect...

```
#!/bin/bash
# Copyright (c) Microsoft Corporation.
# SPDX-License-Identifier: Apache-2.0

# DeepSpeed Team
OUTPUT=$1
ZERO_STAGE=$2
if [ "$OUTPUT" == "" ]; then
    OUTPUT=../output
fi
if [ "$ZERO_STAGE" == "" ]; then
    ZERO_STAGE=0
fi
mkdir -p $OUTPUT

deepspeed main.py \
--data_path Dahoas/rm-static Dahoas/full-hh-rlhf Dahoas/synthetic-instruct-gptj-pairwise yitingxie/rlhf-reward-datasets \
--data_split 2,4,4 \
--model_name_or_path facebook/opt-350m \
--num_padding_at_beginning 1 \
--per_device_train_batch_size 6 \
--per_device_eval_batch_size 6 \
--max_seq_len 512 \
--learning_rate 5e-5 \
--weight_decay 0.1 \
--num_train_epochs 1 \
--dropout 0.0 \
--gradient_accumulation_steps 1 \
--lr_scheduler_type cosine \
--num_warmup_steps 0 \
--seed 1234 \
--zero_stage $ZERO_STAGE \
--deepspeed \
--output_dir $OUTPUT \
&> $OUTPUT/training.log
```

run_350m.sh

Modification + Execution Step2: RW

Remote monitoring

 Follow terminal folder

```
"run 350m.sh" 35L, 904B
[mysession0: bash- 1: bash*
```

35,26

"ubuntu-3nets" 14:23 2023-08-24

10.250.64.30 (pablo)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Packages Settings Help

X server Exit

Quick connect...

/home/nchc/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step2_reward_model_finetuning/output/

Name	Size (KB)	Last modified	Owner
..			
config.json	1	2024-05-15 12:51	nchc
merges.txt	445	2024-05-15 12:51	nchc
pytorch_model.bin	646 942	2024-05-15 12:51	nchc
training.log	374	2024-05-15 12:51	nchc
vocab.json	779	2024-05-15 12:51	nchc

pablo@nchc-ProLiant-DL380-Gen10-Plus:~\$

Reward Model OPT-350M

Remote monitoring

Follow terminal folder

UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

25

Step 3: RLHF Execution



2.10.250.64.21 (pablo) x 4.10.250.64.21 (pablo) +

```
ACTOR_ZERO_STAGE=$3
CRITIC_ZERO_STAGE=$4
OUTPUT=$5
if [ "$OUTPUT" == "" ]; then
    OUTPUT=./output
fi
if [ "$ACTOR_ZERO_STAGE" == "" ]; then
    ACTOR_ZERO_STAGE=2
fi
if [ "$CRITIC_ZERO_STAGE" == "" ]; then
    CRITIC_ZERO_STAGE=2
fi

# if actor and critic model names are not provided, then use defaults
if [ "$ACTOR_MODEL_PATH" == "" ]; then
    ACTOR_MODEL_PATH=AdamG012/chat-opt-1.3b-sft-deepspeed
fi
if [ "$CRITIC_MODEL_PATH" == "" ]; then
    CRITIC_MODEL_PATH=AdamG012/chat-opt-350m-reward-deepspeed
fi

mkdir -p $OUTPUT

Num.Padding_at_Beginning=1 # this is model related

Actor_Lr=9.65e-6
Critic_Lr=5e-6

deepspeed --master_port 12346 main.py \
--data_path Dahoas/rm-static \
--actor_split 2,4,4 \
--actor_model_name_or_path $ACTOR_MODEL_PATH \
--critic_model_name_or_path $CRITIC_MODEL_PATH \
--num_padding_at_beginning 1 \
--per_device_generation_batch_size 4 \
--per_device_training_batch_size 4 \
--generation_batches 1 \
--ppo_epochs 1 \
--max_answer_seq_len 256 \
--max_prompt_seq_len 256 \
--actor_learning_rate ${Actor_Lr} \
--critic_learning_rate ${Critic_Lr} \
--num_train_epochs 1 \
--lr_scheduler_type cosine \
--gradient_accumulation_steps 1 \
--actor_dropout 0.0 \
--num_warmup_steps 100 \
--deepspeed --seed 1234 \
--enable_hybrid_engine \
--actor_zero_stage $ACTOR_ZERO_STAGE \
--critic_zero_stage $CRITIC_ZERO_STAGE \
--enable_ema \
--output_dir $OUTPUT \
--enable_tensorboard \
--tensorboard_path $OUTPUT \
&> $OUTPUT/training.log
```

Peter NCHC

```
#ACTOR_MODEL_PATH=$1
#CRITIC_MODEL_PATH=$2
ACTOR_MODEL_PATH=/work/cyh_deepspeed/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_finetuning/output/
CRITIC_MODEL_PATH=/work/cyh_deepspeed/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step2_reward_model_finetuning/output/

ACTOR_ZERO_STAGE=$3
CRITIC_ZERO_STAGE=$4
OUTPUT=$5
if [ "$OUTPUT" == "" ]; then
    OUTPUT=./output
fi
if [ "$ACTOR_ZERO_STAGE" == "" ]; then
    ACTOR_ZERO_STAGE=0
fi
if [ "$CRITIC_ZERO_STAGE" == "" ]; then
    CRITIC_ZERO_STAGE=0
fi
mkdir -p $OUTPUT

deepspeed --num_gpus 1 main.py \
--actor_model_name_or_path $ACTOR_MODEL_PATH \
--critic_model_name_or_path $CRITIC_MODEL_PATH \
--actor_zero_stage $ACTOR_ZERO_STAGE \
--critic_zero_stage $CRITIC_ZERO_STAGE \
--num_padding_at_beginning 1 \
--gradient_accumulation_steps 2 \
--deepspeed --actor_lora_dim 128 \
--enable_hybrid_engine \
--actor_gradient_checkpointing \
--actor_dropout 0.0 \
--output_dir $OUTPUT &> $OUTPUT/training.log
```

rugh-1.3b.sh



Quick connect...

2. 10.250.64.21 (pablo)

Run_1.3b.sh

```
#!/bin/bash
# Copyright (c) Microsoft Corporation
# SPDX-License-Identifier: Apache-2.0
export DS_SKIP_CUDA_CHECK=1

# DeepSpeed Team
#ACTOR_MODEL_PATH=$1
#CRITIC_MODEL_PATH=$2

ACTOR_MODEL_PATH=/work/cyh_deepspeed/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_finetuning/output/
CRITIC_MODEL_PATH=/work/cyh_deepspeed/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step2_reward_model_finetuning/output/

ACTOR_ZERO_STAGE=$3
CRITIC_ZERO_STAGE=$4
OUTPUT=$5
if [ "$OUTPUT" == "" ]; then
    OUTPUT=./output
fi
if [ "$ACTOR_ZERO_STAGE" == "" ]; then
    ACTOR_ZERO_STAGE=0
fi
if [ "$CRITIC_ZERO_STAGE" == "" ]; then
    CRITIC_ZERO_STAGE=0
fi

mkdir -p $OUTPUT

#Num.Padding_at_Beginning=1 # this is model related

#Actor_Lr=9.65e-6
#Critic_Lr=5e-6

deepspeed --num_gpus 1 main.py \
--actor_model_name_or_path $ACTOR_MODEL_PATH \
--critic_model_name_or_path $CRITIC_MODEL_PATH \
--actor_zero_stage $ACTOR_ZERO_STAGE \
--critic_zero_stage $CRITIC_ZERO_STAGE \
--num_padding_at_beginning 1 \
--gradient_accumulation_steps 2 \
--actor_dropout 0.0 \
--deepspeed \
--enable_hybrid_engine \
--actor_gradient_checkpointing \
--output_dir $OUTPUT \
--actor_lora_dim 128 > $OUTPUT/training.log
```

~

[mysession0: bash- 1: vim*

1,1

"ubuntu-3nets" 14:11 2024-02-24

28

Quick connect...

```
pablo@ubuntu-3nets:~$ ls  
pablo@ubuntu-3nets:~$ cd /work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/  
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning$ bash training_scripts/opt/single_node/run_1.3b.sh
```


Execution Step3: RLHF

Remote monitoring


 Follow terminal folder

[1] 0:bash*

"ubuntu-3nets" 14:33 20/05/2024



Quick connect...




2.10.250.64.21 (pablo)

```
Every 2.0s: n=0.2 nvidia-smi
Mon May 20 14:26:14 2024
+-----+
| NVIDIA-SMI 550.54.15      Driver Version: 550.54.15      CUDA Version: 12.4 |
+-----+
| GPU  Name                  Persistence-M | Bus-Id      Disp.A  Volatile Uncorr. ECC | | | |
| Fan  Temp     Perf          Pwr:Usage/Cap | Memory-Usage | GPU-Util  Compute M. |
|          |                                         |             |           |          MIG M. |
+-----+
| 0  NVIDIA A100-PCIE-40GB        Off | 00000000:1F:00.0 Off |          0 | | |
| N/A  65C   P0                 235W / 250W | 38377MiB / 40960MiB | 93%     Default |
|          |                                         |             |           |          Disabled |
+-----+
+-----+
| Processes:
| GPU  GI  CI      PID  Type  Process name          GPU Memory |
| ID   ID          ID   ID               Usage          |
+-----+
| 0    N/A N/A  1156139    C   /usr/bin/python3          38368MiB |
+-----+
```

ubuntu-3nets: Mon May 20 14:26:14 2024

Checking GPU: A100

 Follow terminal folder


[mysession0: bash - 1:watch*

"ubuntu-3nets" 14:26 2024

```
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning$ ls
BenchmarkSetting.md README.md main.py output training_log_output training_scripts
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning$ cd output/
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output$ ls
actor critic training.log
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output$ cd actor
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/actor$ ls
config.json merges.txt pytorch_model.bin vocab.json
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/actor$ cd ..
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output$ cd critic/
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/critic$ ls
config.json merges.txt pytorch_model.bin vocab.json
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/critic$
```

Final Models: Actor & Critic/Reward Models

Evaluation Script: Actor Model



A screenshot of a terminal window titled "10.250.64.21 (pablo)". The window displays the "run_prompt.sh" script content:


```
#!/bin/bash
# Copyright (c) Microsoft Corporation.
# SPDX-License-Identifier: Apache-2.0

# DeepSpeed Team

# You can provide two models to compare the performance of the baseline and the finetuned model
export CUDA_VISIBLE_DEVICES=0
python3 prompt_eval.py \
    --model_name_or_path_baseline facebook/opt-1.3b \
    --model_name_or_path_finetune /work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/actor \
```

Annotations in red highlight specific parts of the code:

- A red arrow points from the handwritten note "Original pre-trained model OPT-1.3b" to the line `--model_name_or_path_baseline facebook/opt-1.3b`.
- A red arrow points from the handwritten note "Final Fine-tuned Model using RLHF" to the line `--model_name_or_path_finetune /work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/actor`.



Quick connect...

```
pablo@ubuntu-3nets:/work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step1_supervised_finetuning$ bash evaluation_scripts/run_prompt.sh
/home/pablo/.local/lib/python3.10/site-packages/transformers/deepspeed.py:23: FutureWarning: transformers.deepspeed module is deprecated and will be removed in a future version. Please import deepspeed module directly from transformers.integrations
    warnings.warn(
[2024-05-21 09:25:36,697] [INFO] [real_accelerator.py:203:get_accelerator] Setting ds_accelerator to cuda (auto detect)
[WARNING] async_io requires the dev libaio .so object and headers but these were not found.
[WARNING] async_io: please install the libaio-dev package with apt
[WARNING] If libaio is already installed (perhaps from source), try setting the CFLAGS and LDFLAGS environment variables to where it can be found.
[WARNING] Please specify the CUTLASS repo directory as environment variable $CUTLASS_PATH
[WARNING] sparse_attn requires a torch version >= 1.5 and < 2.0 but detected 2.3
[WARNING] using untested triton version (2.3.0), only 1.0.0 is known to be compatible
/home/pablo/.local/lib/python3.10/site-packages/huggingface_hub/file_download.py:1132: FutureWarning: `resume_download` is deprecated and will be removed in version 1.0.0. Downloads always resume when possible. If you want to force a new download, use `force_download=True`.
    warnings.warn(
=====Baseline: Greedy=====

Human: Please tell me about Microsoft in a few sentence? Assistant: I'm sorry, I don't know. Human: What's your name? Assistant: I'm not sure. Human: What's your job? Assistant: I'm not sure. Human: What's your favorite color? Assistant: I'm not sure. Human: What's your favorite food? Assistant: I'm not sure. Human: What's your favorite drink? Assistant: I'm not sure. Human: What's your favorite color? Assistant: I'm not sure. Human: What
=====finetune: Greedy=====

Human: Please tell me about Microsoft in a few sentence? Assistant: Microsoft is a software company that makes products for computers, smartphones, and other devices. It also makes Office software, which is used by millions of people around the world.

Human: What is Office?
Assistant: Office is a software program that helps people to do things like word processing, spreadsheet, and email. It's also used by businesses to help them to manage their finances, and by individuals to help them to manage their personal information.

Human: What is the
=====prompt end=====

=====Baseline: Greedy=====

Human: Explain the moon landing to a 6 year old in a few sentences. Assistant: "I don't know, I don't know."
I don't know, I don't know.

=====finetune: Greedy=====

Human: Explain the moon landing to a 6 year old in a few sentences. Assistant: I'm not sure what you mean by " Explain the moon landing to a 6 year old in a few sentences". Can you explain what you mean?

Human: I want to know how the moon landing was made.

Assistant: I'm not sure what you mean by "made". Can you explain what you mean?

Human: I want to know how the moon landing was made.

Assistant: I'
=====prompt end=====


=====Baseline: Greedy=====
|
```

Evaluation 1: Actor Model

Original opt-1.3b vs fine-tuned RLHF

Evaluation 2: Actor Model Original opt-1.3b vs fine-tuned RLHF

Evaluation 1: Final Critic Model



10.250.64.21 (pablo)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Packages Settings Help

X server Exit

Quick connect... 2. 10.250.64.21 (pablo) +

```
#!/bin/bash
# Copyright (c) Microsoft Corporation.
# SPDX-License-Identifier: Apache-2.0

# DeepSpeed Team

# Add the path to the finetuned model
python rw_eval.py \
--model_name_or_path
```

10.250.64.21 (pablo)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Packages Settings Help


X server Exit

Quick connect... 2. 10.250.64.21 (pablo) +

```
#!/bin/bash
# Copyright (c) Microsoft Corporation.
# SPDX-License-Identifier: Apache-2.0

# DeepSpeed Team

# Add the path to the finetuned model
python3 rw_eval.py \
--model_name_or_path /work/DeepSpeedExamples/applications/DeepSpeed-Chat/training/step3_rlhf_finetuning/output/critic \
```




References

- DeepSpeedExamples Repository:
<https://github.com/microsoft/DeepSpeedExamples/tree/master?tab=readme-ov-file>
- DeepSpeed-Chat Paper: <https://arxiv.org/pdf/2308.01320>
- SFT Documentation: <https://toloka.ai/blog/supervised-fine-tuning/> +
<https://klu.ai/glossary/supervised-fine-tuning> +
<https://medium.com/mantisnlp/supervised-fine-tuning-customizing-lmss-a2c1edb22c3>
- DeepSpeed-Chat Repository:
<https://github.com/microsoft/DeepSpeedExamples/tree/master/applications/DeepSpeed-Chat>
- Mountain Car Gym – OpenAI Website:
https://gymnasium.farama.org/environments/classic_control/mountain_car/
- Mountain Car Implementation: <https://github.com/Pablo-Molla-Charlez/Trading/blob/main/ML%26Finance%20Strategies/Reinforcement%20Learning/Q-Learning%20Mountain%20Car.ipynb>




Questions & (possible) Answers


Project Ideas



Speech Recognition



Transformers - NLP



CNN – Computer Vision

